

Hybrid Rule-Based Deep Q-Network Agent for LUX AI

Algorithm 1 Overall Agent Architecture

```
1: Initialize Components:
2: Create Q-network (neural network for action evaluation)
3: Create target network (stable copy of Q-network)
4: Create replay buffer (stores past experiences)
5: Set hyperparameters: learning rate, discount factor, exploration rates
6: Initialize tracking systems: relic memory, productive tiles, unit assignments
7:
8: for each game do
9:   Reset game-specific memory
10:  for each match (3 matches per game) do
11:    for each time step do
12:      Observe current state and update relic positions
13:      for each active unit do
14:        Generate random number  $u$  from [0,1]
15:        if  $u <$  DQN probability then
16:          Select action using DQN policy (Algorithm 3)
17:        else
18:          Select action using rule-based policy (Algorithm 4)
19:        end if
20:        Augment with attack action if appropriate (Algorithm 5)
21:      end for
22:      Execute all unit actions
23:      Observe rewards and next state
24:      Update productive tile statistics (Algorithm 6)
25:      Calculate shaped rewards for each unit (Algorithm 7)
26:      Store transitions in replay buffer
27:      if training step then
28:        Perform network update (Algorithm 8)
29:      end if
30:    end for
31:    Reset match-specific memory
32:  end for
33:  Reset all memories for next game
34: end for
```

Algorithm 2 Neural Network Architecture

- 1: **Q-Network Structure:**
- 2: Input layer: Game state features (position, energy, distances, etc.)
- 3: Hidden layer 1: 256 neurons with ReLU activation
- 4: Hidden layer 2: 256 neurons with ReLU activation
- 5: Hidden layer 3: 256 neurons with ReLU activation
- 6: Hidden layer 4: 128 neurons with ReLU activation
- 7: Output layer: 5 neurons (Q-values for each action)
- 8: Actions: 0=Stay, 1=Up, 2=Right, 3=Down, 4=Left
- 9:
- 10: **Target Network:**
- 11: Same architecture as Q-network
- 12: Updated periodically (every N training steps) for learning stability

Algorithm 3 DQN Action Selection (Epsilon-Greedy)

Require: Current state, Q-network, exploration rate

- 1:
- 2: Generate random number u from [0,1]
- 3: **if** $u <$ exploration rate **then**
- 4: **return** random action *// Exploration*
- 5: **else**
- 6: Compute Q-values for all actions using Q-network
- 7: **return** action with highest Q-value *// Exploitation*
- 8: **end if**

Algorithm 4 Rule-Based Policy

Require: Unit ID, time step, game state, relic positions, productive tiles

```
1:  
2: Extract unit position, energy level, and active status  
3: if unit inactive OR energy below safety threshold then  
4:   return Stay action  
5: end if  
6:  
7: Determine number of active units  
8: Calculate unit's rank among active units  
9: Compute exploration ratio based on game phase:  
10: Early matches: higher exploration (find new relics)  
11: Late matches: lower exploration (focus on collection)  
12:  
13: Assign unit role based on rank and exploration ratio  
14:  
15: if unit is COLLECTOR then  
16:   Priority 1: Stay on productive tiles  
17:   if current position is productive tile then  
18:     return Stay action  
19:   end if  
20:  
21:   Priority 2: Move to known productive tiles  
22:   if productive tiles exist then  
23:     Find nearest productive tile  
24:     return direction toward nearest productive tile  
25:   end if  
26:  
27:   Priority 3: Systematic grid exploration  
28:   if unit has no assignment then  
29:     Assign unit to position in 5×5 grid around relics  
30:   end if  
31:   Calculate target position from assignment  
32:   return direction toward assigned grid position  
33: else  
34:   // Unit is EXPLORER  
35:   Systematic quadrant search  
36:   if unit has no exploration target then  
37:     Assign quadrant based on unit ID (4 quadrants total)  
38:     Generate random target within assigned quadrant  
39:   end if  
40:  
41:   if target reached OR stuck at current target then  
42:     Generate new random target in quadrant  
43:   end if  
44:   return direction toward exploration target  
45: end if
```

Algorithm 5 Attack Action Selection

Require: Unit position, game state, unit energy, attack range

```
1:  
2: Get opponent positions and ally positions  
3: if energy insufficient OR no opponents visible then  
4:   return no attack  
5: end if  
6:  
7: Initialize best score =  $-\infty$  and best target = none  
8: for each possible target location within attack range do  
9:   Calculate direct hit score (enemies at exact target)  
10:  Calculate splash score (enemies adjacent to target)  
11:  Calculate friendly fire penalty (allies at or near target)  
12:  Calculate distance penalty (slight preference for closer targets)  
13:  
14:  Total score = direct + splash - friendly fire - distance  
15:  
16:  if score > best score then  
17:    Update best score and best target  
18:  end if  
19: end for  
20:  
21: if best score  $\geq$  minimum threshold then  
22:   return attack toward best target  
23: else  
24:   return no attack  
25: end if
```

Algorithm 6 Productive Tile Detection

Require: Current state, team points, previous points, position statistics

```
1:  
2: Calculate points gained this step = current points – previous points  
3: Find all unit positions near known relics (within distance 3)  
4:  
5: for each position near relics do  
6:   Increment occupancy counter for this position  
7: end for  
8:  
9: if points gained > 0 AND units are near relics then  
10:   Distribute gained points among near-relic positions  
11:   for each position that was occupied do  
12:     Update total points earned at this position  
13:     Calculate confidence = total points / times occupied  
14:  
15:     if confidence > 0.5 AND occupied at least 3 times then  
16:       Mark position as "productive tile"  
17:     end if  
18:   end for  
19: end if  
20:  
21: Periodic cleanup (every 100 steps):  
22: Remove tiles from productive set if confidence drops below 0.7
```

Algorithm 7 Reward Shaping Function

Require: State, action, team reward, productive tiles, statistics

```
1:  
2: Initialize shaped reward =  $3.0 \times$  team reward  
3: Extract current position and next position from states  
4:  
5: Component 1: Productive Tile Bonuses  
6: if currently on productive tile then  
7:   Add large bonus scaled by tile confidence  
8:   if stayed on same tile (didn't move) then  
9:     Add additional staying bonus  
10:  end if  
11: end if  
12: if moved to a productive tile then  
13:   Add bonus for reaching productive tile  
14: end if  
15:  
16: Component 2: Distance Improvement  
17: if productive tiles are known then  
18:   Calculate distance improvement toward nearest productive tile  
19:   Add reward proportional to improvement  
20: else  
21:   Calculate distance improvement toward nearest relic  
22:   Add smaller reward proportional to improvement  
23: end if  
24:  
25: Component 3: Energy Management  
26: Calculate energy change  
27: if energy increased then  
28:   Add small bonus for energy recovery  
29: end if  
30: if energy below threshold AND not on productive tile then  
31:   Add penalty for risky low energy state  
32: end if  
33:  
34: Component 4: Movement Encouragement  
35: if unit moved AND not leaving productive tile then  
36:   Add small bonus to encourage exploration  
37: end if  
38:  
39: Component 5: Survival Bonus  
40: Add small constant bonus for staying alive  
41:  
42: Component 6: Leaving Productive Tile Penalty  
43: if left productive tile without going to another productive tile then  
44:   Add penalty scaled by confidence of abandoned tile  
45: end if  
46:  
47: return shaped reward
```

Algorithm 8 Training Procedure (Double DQN)

Require: Replay buffer, batch size, Q-network, target network

```
1:  
2: if not enough experiences in buffer then  
3:   return // Skip training  
4: end if  
5:  
6: Sample random batch of experiences from replay buffer  
7: Each experience contains: (state, action, reward, next state, done flag)  
8:  
9: Step 1: Compute current Q-value predictions  
10: Pass states through Q-network  
11: Extract Q-values for actions that were actually taken  
12:  
13: Step 2: Compute target Q-values (Double DQN method)  
14: Use Q-network to SELECT best actions for next states  
15: Use target network to EVALUATE those actions  
16: Calculate targets = reward + discount × evaluated Q-value  
17: If episode ended, target = reward only (no future value)  
18:  
19: Step 3: Calculate loss and update  
20: Compute mean squared error between predictions and targets  
21: Calculate gradients of loss with respect to network parameters  
22: Clip gradients to prevent instability  
23: Update network parameters using Adam optimizer  
24:  
25: Step 4: Decay exploration rate  
26: Reduce exploration rate: new rate = current rate × decay factor  
27: Ensure rate doesn't go below minimum value  
28:  
29: Step 5: Periodically update target network  
30: if training step is multiple of update frequency then  
31:   Copy Q-network parameters to target network  
32: end if
```

Key Concepts Explained

Hybrid Policy

The agent combines two decision-making strategies:

- **Rule-based policy:** Hand-crafted rules for reliable behavior
- **DQN policy:** Learned through experience via neural networks
- A probability parameter controls which policy is used

Q-Network

A neural network that estimates the "quality" (Q-value) of each action in a given state. Higher Q-values indicate better expected future rewards.

Experience Replay

Past experiences are stored in a buffer and randomly sampled during training. This breaks correlations between consecutive experiences and improves learning stability.

Target Network

A slowly-updated copy of the Q-network used for computing learning targets. This prevents the "moving target" problem and stabilizes training.

Reward Shaping

The basic team reward is augmented with additional signals to guide learning:

- Bonuses for productive tiles
- Distance improvement rewards
- Energy management signals
- Exploration encouragement

Productive Tile Detection

Statistical tracking identifies which positions near relics actually generate points. Confidence increases as more evidence accumulates.